

Climate control for performant HPUs

April 10, 2014

Abstract

Introduction

Although capabilities in artificial intelligence continue to advance, there are still many tasks for which human performance exceeds that of computers now and in the foreseeable future. Tasks requiring repertoire of general knowledge about the world, the use of common sense or expert judgment, finding creative solutions to open-ended questions, and spontaneously hypothesis generation outside any apparent framework, are all examples of tasks that are routine for humans, but very difficult or impossible for computers. The image labelling task is perhaps the canonical human intelligence task (HIT), perhaps because it is simply stated, nearly effortless for humans, but still extremely challenging for computers.

Rather than replicate human intelligence, there is an alternate vision in which human intelligence can be freely accessed and meshed with artificial intelligence. The emergence of micro-task platforms like Amazon Mechanical Turk (AMT) has made this vision a reality to some extent. In this paradigm, the human processing unit (HPU) becomes one component in a hybrid computational system, analogous to the CPU.

But there are major challenges in building systems that deliver on this vision. People are far more complex than manufactured chips, and

so HPU performance is noisy, and subject to bias, leading to uncertain quality of output. Of course, the effect that HPU noise has on a compute job will depend on the application. In the image labelling task, some variability is desirable, because it helps generate sets of labels that in some sense cover semantic space occupied by an image. In other cases this variability can be problematic. In a transcription task, usually there is only one correct output, and so variability in the output only degrades quality.

Regardless of the application, the understanding HPU variance quantitatively, and the factors that influence it, will help design of efficient compute systems built from HPUs.

Much has been written on the factors that influence human input in crowdsourcing platforms. This has generally related to the context in which the HPU engages the task. Unfortunately, especially for platforms like AMT, this context is out of the designers control, and the relative anonymity of workers makes it difficult to select for specific conditions.

In the present work, we compare influences that arise from the framing of the task to influences that arise in-task from the mere act of working. What are the effects of performance that arise simply from how the worker's attention is influenced by the previous task, and how do these relate quantitatively to other pre-task influences?

We pay workers on AMT to label images, and analyze the effects brought about by subjecting them to different priming treatments, either in the form of disclosing a semantically charged name of a (fictitious) organization funding the

research, or by altering the first few images presented to the workers. Surprisingly, simply altering the first few images produces a much stronger shift in subsequent labelling performance than does disclosing the semantically-laden name of the funder.

Prior Work

When considering the design of computing systems based on HPUs, we can think of variability in HPU output as having a relatively persistent component, and a transient one. The persistent component would include intrinsic qualities of a person, such as their temperament, life history, and their current developmental stage. Such characteristics cannot be influenced by the designer, but it may be possible to screen or at least characterise them to some extent.

The transient component of variability might include such factors as influence alertness, orientation focus, mood, meaningfulness of the task, and any number of aspects of mental state too subtle to describe. Such aspects are partly in the control of the system designer, and can be controlled through how the task is framed and set up.

In psychology, the phenomenon of priming occurs when a prior stimulus makes a person more likely to respond to a task in a particular way. The phenomenon is well-studied, and is believed to follow perceptual, semantic, or conceptual similarities between the priming stimulus and the ensuing response.

Drawing from this insight, it is natural to wonder whether current HPU performance might be modulated by recently performed tasks.

Existing work on the effects of human input in crowdsourcing platforms has generally focused on effects arising from auxilliary information, which has been added into the task setting, or which frames the task during its introduction.

For example, it was shown that framing a task

in a meaningful or meaningless way influences both task quality and the willingness of workers to produce more output given a declining payment schedule [1]. Surprisingly, relative to a zero-context treatment, although workers who were told that their contributions would be used to help identify cancerous cells were willing to provide more work, their work was not of a sufficiently higher quality. The only effect on quality occurred when workers were told that their output would be discarded, in which case quality declined.

Another study investigated the influence of providing workers with information about one another. This study found that workers were willing to provide more output, when they were able to see one another’s names. The effect was stronger when workers could also see some basic self-reported demographic information, and more still when they could see one another’s responses.

In the present work, we focus on a perhaps unusual source of priming—that arising from the task itself, and compare this to priming that results from framing the work as part of study by a named organization. Given that the mechanisms of priming are believed to be related to the residual activation of perceptual, semantic, and conceptual representations, we hypothesize that *in-task* priming could produce effects quantitatively comparable to those produced by framing.

Theoretical Framework

Having discussed very briefly some mechanisms and cases of priming, we shall now seek to put forward, as an additional testable contribution, a rigorous definition for priming that is suitable for the study of computing with HPUs.

First, we remark that it does not make sense to speak of an un-primed state. When a person engages in a task, she comes to the task with some state of mind. We therefore do not attempt to

define priming in an absolute sense, but rather, in a relative sense, that is, one priming can be said to be different from another.

Also, the effects of a given treatment will depend on the task to which the HPUs are put. In other words, a treatment which generates a strong effect on the output in one task, does not necessarily produce a strong effect on the output for other tasks. On a related note, the study of the effects of priming on HPUs should be defined algorithmically. Only effects on HPU output that can be detected algorithmically can be relevant to a computing system built from HPUs.

Further, since only the influence on populations of HPU outputs can be meaningfully measured, and since accessing HPU power generally involves sampling from a pool of HPU workers, we define priming as a property of a population of HPU outputs. Having said these remarks, we now present an algorithmic definition of HPU priming:

Two populations of HPU outputs, J and K , are said to be *differently primed* with respect to a task \mathcal{T} if there exists an algorithm \mathcal{A} which runs in time polynomial in the size of J and K , that can distinguish (classify) members of J and K with accuracy $\frac{1+\theta}{2}$, on input J and K . Further, θ must be a non-negligible function of $|J + K|$. If such an \mathcal{A} exists, we say that J and K deviate by θ in priming.

The above definition is simply intended to provide a well-defined definition of what priming *is*. Naturally, it says nothing about the consequences of priming. The significance of the priming of a given population will depend both on the nature of \mathcal{T} and on the intended purpose of the work products derived from HPUs performing \mathcal{T} .

Methods

Task set-up. We paid 900 AMT workers to perform an image-labelling task. A task con-

sisted of labelling 10 images, with 5 labels each. The first 5 images were varied depending on the priming treatment, while the last 5 images were the same across all treatments. Ordering of the images was kept constant.

Workers were randomly assigned to one of 6 treatments. The treatments differed from one another along two dimensions. The first dimension consisted of varying the first 5 images shown to the worker. This was used to test the effects of *in-task* priming.

The second dimension concerned disclosure of a (fictitious) organization, purportedly funding work as part of a research study. Depending on the treatment, one of two funding agencies was presented, or no indication was made.

Tasks were presented to workers as a series of panels or flash cards. The first panel provided brief instructions, and was identical for all treatments. Workers could see this panel when previewing the task, but could not advance. Depending on the treatment, the worker was either shown a second panel stating the name of one of two fictitious organizations funding the work, or this panel was skipped. The next five panels each consisting of a priming sub-tasks, wherein the worker was asked to submit five descriptive labels. The images used during the priming sub-task depended on the treatment. The last 5 panels consisted of testing subtasks, wherein, as for the priming sub-tasks, workers were asked to submit 5 descriptive labels.

Choice of images. The 5 test images, were chosen with two ideals in mind. First, we chose images that we judged would generate a diverse vocabulary of labels, such that the effects of priming could be detected. In other words, sparse images with a single object in the foreground were not considered good candidates, since they would be less likely to elicit labels that varied from one worker and one priming treatment to the next.

Second, we chose images which would produce labels belonging to two broad concepts, which would serve as the targets of our priming:

food and culture. This created the opportunity to attempt to prime workers in a way that would bias them toward emitting food-related or culture-related labels.

Under these considerations, we chose the images shown in Fig. 5. Each of these images has food as its main focus, but also has a strong and specific cultural reference due to the unique, iconic character of the food and the artifacts depicted.

To investigate in-task priming, we chose a set of images that highly recognizeable cultural settings and no food, and another set that contained separated food ingredients, without any overt cultural content. The third set of images was chosen to be very much like the test images, showing prepared meals, and though prepared food is inseparable from culture, these images were chosen based on being culturally more muted or ambiguous.

Label ontology. In order to provide a deeper analysis, we built an ontology of the corpus of all labels applied to the first test image. The ontology was built as a directed acyclic graph starting

Results and discussion

Priming affects HPU output. Before looking for differences in the content of labels provided by workers from different treatments, we first demonstrate that the treatments are distinguishable in an algorithmic sense.

Using a naive bayes classifier, we are able to distinguish with high precision and specificity between workers from the AMBG treatment and any of the other 5 treatments. Fig 1A shows F1-score for a naive bayes classifier when distinguishing between the AMBG and the other treatments. The classifier achieves high precision and specificity in task. Figure 1B shows the F1-score for binary classification between the $CULT_{img}$ and the other treatments, while 1C shows that for $INGR_{img}$ and the other treat-

ments.

Not surprisingly, pairs of treatments in which one primes for culture, and the other primes for ingredients are easily distinguished. However, it is quite surprising that high accuracy is achieved when classifying treatments within the same orientation (e.g. cultural). This is especially true in such cases as the classification of $CULT_{img}$ and $CULT_{fund,img}$ ($F_1 = 0.863$), where both treatments also share the priming images in common.

We find it remarkable that, using only the labels that workers provide, it is possible to infer with good accuracy (at least when given a choice between two possibilities) the treatment to which the worker was subjected.

Priming orients HPU focus. Having established that each treatment does in fact consist of distinguishable populations of the HPUs, we next look at the content of labels. Two of the root labels in our ontology, **food** and **cultural**, map naturally onto the two concepts that laid behind the design of our priming treatments. A natural and straightforward expectation is that HPUs from treatments $CULT_x$ should emit more labels that are ontological descendents of **cultural**, while those from the $INGR_x$ treatments should emit more labels under **food**.

In Fig. 2, we exhibit the percentage of labels having a cultural orientation (i.e. descending from **cultural** in the ontology), and the percentage having a food orientation. Before making comparisons using this information, we remark that there is no reason to expect a balance in the number of words having cultural or food orientation overall. The AMBG treatment, which was designed to promote neither orientation, exhibits a significant fraction of words from both the food and cultural orientation, while significantly favoring the former.

Both $CULT_{img}$ and $CULT_{fund,img}$ show a significant excess of culturally-oriented labels and fewer food-oriented labels. Furthermore, HPUs in these treatments emit more culturally-oriented labels even while emitting fewer labels

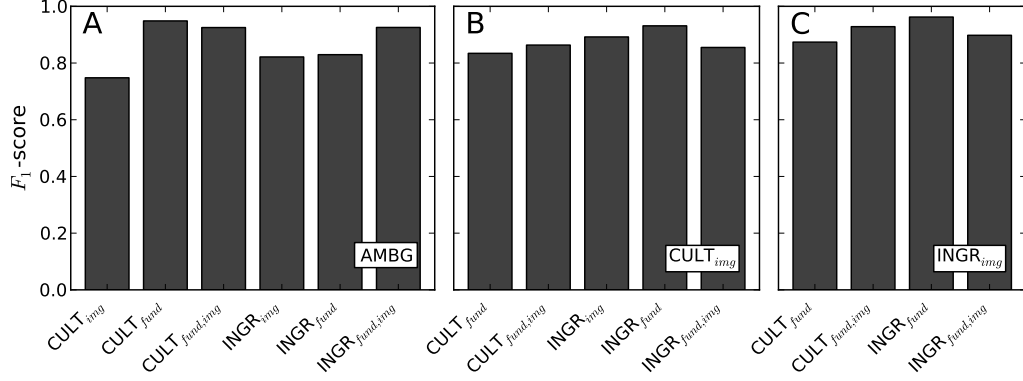


Figure 1: F_1 -score for binary classification of HPUs from separate treatments using a naive Bayes classifier. Each pannel shows the performance of the classifier when distinguishing between a basis treatment (inset) and the treatments listed on the abscissa.




Figure 2: Percentage of labels of a food- or cultural-orientation, or both. In our ontology of labels, a label can have multiple parents. For example **naan** inherits from both **food** and **cultural** through its parent **indian food**.

that are simultaneously oriented toward cultural and food (light bars in Fig. 2). The deviation of these treatments from the others is well beyond the 95% confidence interval.

Interestingly, the $CULT_{fund}$ treatment did not have this effect. Although we have demonstrated that $CULT_{fund}$ is differently primed from AMBG using a naive bayes classifier, in respect of overall fraction of food- and culturally-oriented labels, no distinction is to be made.

In the case of $INGR_x$ treatments one expects to see an enrichment of food-oriented labels. There is perhaps some evidence for this $INGR_{fund}$ and $INGR_{fund,img}$, but we cannot make any assertion with confidence.

Priming affects attention to detail. We next look at how alternately treated HPUs differ in their tendency to use more specific or more general labels. We use the ontology of labels emitted on the first test image to unambiguously establish a partial ordering of label-specificity. If one label ℓ_1 is within the ancestry of another ℓ_2 , we say that ℓ_1 is more general than ℓ_2 , otherwise they are not comparable. Thus, the label **naan** is more specific than both **bread** and **indian**, while uncomparable to **statue**.

Label-specificity only generates a partial

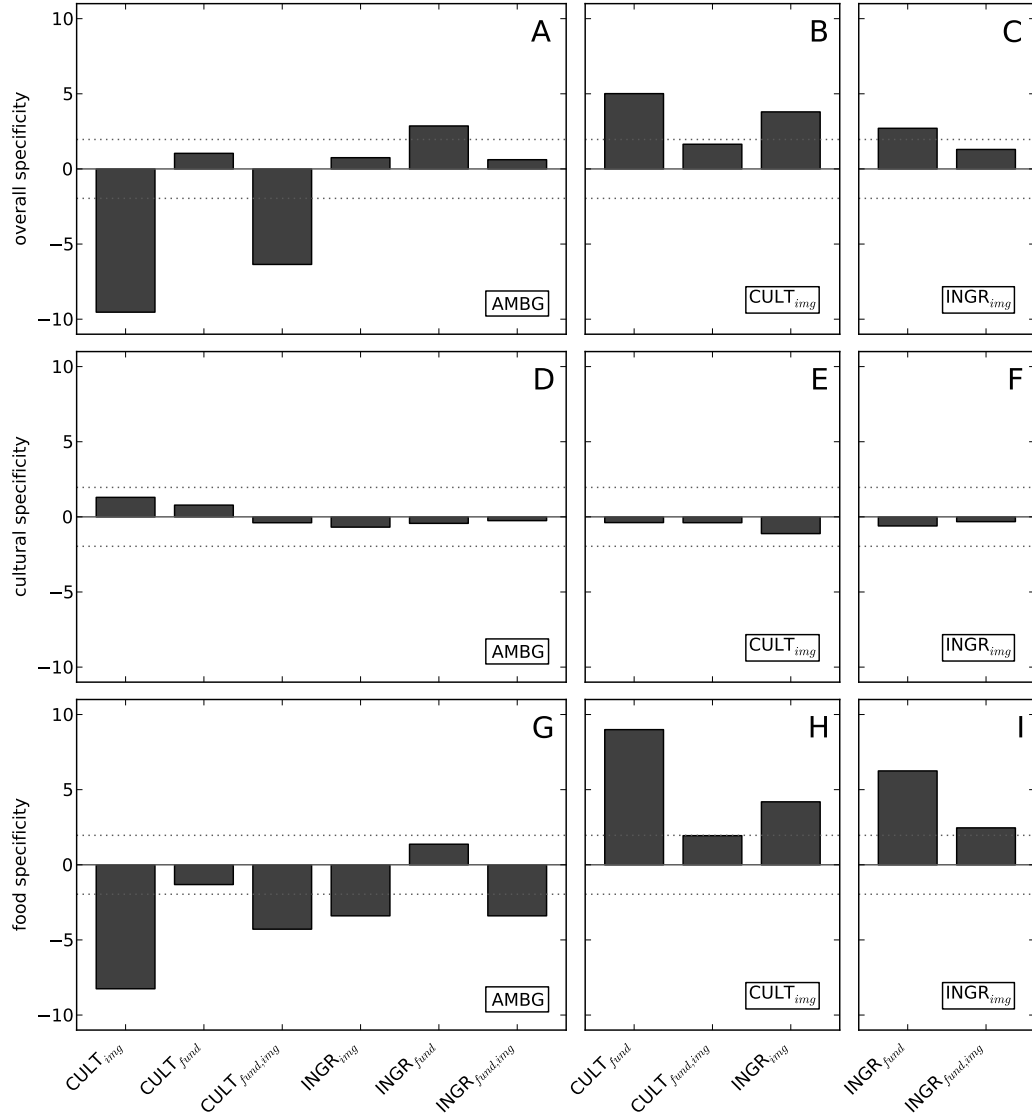


Figure 3: Pairwise comparisons of label specificity between different HPU treatments. Each panel presents a binary comparison between a basis treatment (inset) and the subject treatments indicated on the abscissa. A positive specificity score indicates that the subject treatment emitted more specific words than the basis treatment overall. In a given comparison, a sample of 50 HPUs from the each treatment was randomly sampled, and the specificity of labels from the HPUs of opposing treatments were compared. To compare two HPUs, each pair of labels from different HPUs are compared, and the specificity score of subject HPU is the number of cases where its label is more specific than the subject HPUs, less the number of cases where it is more general. This score is averaged for all HPU pairings. Statistical significance is gauged by generating a null-comparison between two mutually exclusive subsamples from the basis treatment. This null-comparison yields a distribution of relative specificities whose mean is in principle zero. The specificity scores are expressed in terms of standard deviations of the null-comparison specificities. The dotted lines represent the 95% confidence interval for rejecting the null Hypothesis that the basis treatment and subject treatment are equally specific.

ordering—at least in our experimental design. A consequence of this is that it is not possible to assign an overall specificity score to a set of labels. We submit that this reflects the underlying complexity of natural language semantics. While we admit that there may be many ways to generate full orderings based on some notion of label specificity, we believe that this would inappropriately collapse qualitative differences between labels, leading to results that are difficult to interpret.

In Fig. 3, We show pairwise specificity comparisons between various pairs of treatments. Figure 3A shows the comparison of AMBG with all other treatments. Both treatments primed with cultural images ($CULT_{img}$ and $CULT_{fund,img}$) exhibit an excess of more general words compared to AMBG. Meanwhile, $INGR_{fund}$ emitted an excess of more specific words.

We can seek to explain these observations by restricting the comparison to labels of a specific orientation (e.g. food or cultural). When we perform the same comparisons but restricting to food-oriented labels, as shown in Fig. 3G, we see some of the same tendencies as in A. However, when restricting to culturally-oriented labels (Fig 3D) there are no significant comparative deviations in specificity to speak of.

A first conclusion that we can draw from this is that there is a significant loss of specificity in labels emitted by treatments having non-ambiguous priming images. In Fig. 3G, only $CULT_{fund}$ and $INGR_{fund}$ show no significant deviation.

- through negative priming workers stop responding to generic features of images of prepared meals, which happens in the ambiguous case.

- On the other hand, workers presented with either the cultural or ingredients images are initially struck by the gross differences, which draws generic labels, especially with respect to food for which generic labels in ambiguous treatment had been suppressed through negative priming.

- it seems, at first, inconsistent that $CULT_{img}$ and $CULT_{fund,img}$ would at once enrich the fraction of culturally-oriented words, but produce no change in the specificity of culturally oriented words. First of all, there is no logical incompatibility with producing a greater number of cultural terms, yet which are generic in nature. If we follow the hypothesis that focus (or orientation) is directed by positive priming (perhaps through the subconscious or conscious inference of requester intent), while focusing on nuances comes from the inhibition of gross features through negative priming, then this combination of observations makes sense. Having seen many cultural images, HPUs in $CULT_{img}$ and $CULT_{fund,img}$ are primed to emit cultural labels, however, since the diversity of images is high compared to other treatments on reaching the test images, there no opportunity for focussing in to finer detail through the mechanism of negative priming.

In-task priming is stronger than framing.

Our experimental set-up includes two priming mechanisms: in-task priming, produced by varying the first 5 images of the task, as well as what could be called sidestream priming, in which a fictitious funder is disclosed. The intention behind this set-up was to enable a direct comparison, and our expectation was that in-task priming might produce some fraction of the effect produced by disclosing the funder. To our surprise, in-task priming produces a much stronger effect.

There is a cautionary lesson here. In-task priming, which is inherently impossible to eliminate is may be far more severe than the more overt causes of priming that more routinely attract concern during the design of an experiment. Depending on the final purpos of HPU work products, this may be quite restrictive.

Leveraging priming. In a more positive view, the results of our study suggest that the ordering of subtasks can be used as a way to direct the focus and attention to detail by HPUs. If the designer desires more coarse-grained work-products, this can be achieved by presenting




Figure 4: F_1 -score for a naive Bayes classifier trained to distinguish HPUs from $INGR_{img}$ and $CULT_{img}$ based on 80 instances from each treatment, and tested using 20 instances. Here we present the performance of the classifier when only the labels from specific images are presented, to exhibit whether its performance degrades as the number of subtasks since priming increases.

successive subtasks with high diversity. On the other hand if fine-grained work products are desired, then the designer should aim to sort subtasks into “tracks” that are highly homogeneous, and assign subsens of the pool of HPUs to specifict tracks.

It is interesting to consider to what extent the output of HPUs can be driven toward nuanced detail using this technique. Consider, for example, the image labelling platform called the ESP Game. Here Two HPUs are shown the same image and in order to derive what are in some sense the moset characteristic labels, they are asked to attempt to produce the same labels for the image. One could imagine a similar setup, but in which the designer attempts to produce the set of labels that best covers the semantic space occupied by the image, by eliciting labels at coarser and finer levels of specificity using the techniques suggested by the present work.

One interesting test of this hypothesis arises in

References

- [1] Dana Chandler and Adam Kapelner. Breaking monotony with meaning: Motivation in crowdsourcing markets. *Journal of Economic Behavior & Organization*, 90:123–133, 2013.

Figure 5: caption here

Figure 6: caption here

Figure 7: caption here

Figure 8: caption here