

# Climate control for performant HPUs

April 9, 2014

## Abstract

## Introduction

There are still many tasks for which humans outperform computers. These tend to be tasks that rely on a large repertoire of prior knowledge, the exercise of common sense judgment, certain visual tasks, and tasks that require spontaneous hypothesis generation.

With the emergence of microtask platforms like Amazon Mechanical Turk (AMT), comes the hope of building compute in which human intelligence can be accessed on-demand. This leads to an analogy of the human to the CPU, wherein we think of human processing units.

A major challenge in building systems that deliver on this vision is the natural variability obtained in HPU output. In some tasks, this variability is desirable. For example, in an image labelling task, the natural variability in HPU output leads to greater coverage of the semantic space occupied by an image. In other cases this variability can be problematic, such as in a transcription task, where there is only one correct output, and any variability is only noise.

Whether desired or not, to build up efficient systems from HPUs, one must understand the natural variance quantitatively, and understand the factors that influence it.

Here we model the variance between HPUs as coming from two sources. The first is a per-

sistent, or intrinsic variability. We imagine that this to encompass temperamental predispositions, life history, and developmental stage of the individual. While it may change over the course of a person's lifetime, in the view of the computational architect, it is essentially constant.

The second is a short-term variability, which arises from the fact that people's immediate state is a function of recent events. This would include such things as an HPU becoming more efficient at a task once it has completed a few tasks of the same type. It would also include such things as the potentially detrimental effects of a noisy environment.

In this view, one can imagine HPUs being different in their specifications, and also exhibiting hysteresis, whereby their output at some time  $t$  is a function of their inputs at time  $t$ ,  $t - 1$ , and so on.

In psychology, the act of producing these short-lived states that effect a person's performance during a task is called *priming*. We will adopt this term throughout the present work.

## Prior Work

Priming itself can be broken into many types, and many studies exist exploring the effects of overt priming on the performance of HPUs. Priming may arise due to how the task is framed. This includes such things as the stated purpose of the task, and the identity of the requester the task.

In [1], the researchers investigate the effects of framing task either in a meaningful or meaningless way. Compared to a zero-context control treatment, workers increase their output (for less pay), when they are told that they are helping identify cancer, but there is no change in quality. When workers are told that their submissions will be discarded, there is no change in the amount of work done, but the quality declines.

Another source of priming can be due to what might be called *sidestream information*: information that is presented to the individual but which is not actually salient to the task.

Other researchers have studied how having peer's responses available to the task could influence the worker.

In the present work, we focus on a more subtle source of priming, which arises simply from the worker's performance in earlier stages of the task.

We compare this to priming arising from framing the task by disclosing a (fictitious) funding agency that is paying for the study in which they are participating.

## Methods

We paid 900 AMT workers to perform an image-labelling task. A task consisted of labelling 10 images, with 5 labels each. The first 5 images were varied depending on the priming treatment, while the last 5 images were the same across all treatments. Ordering of the images was kept constant.

Workers were randomly assigned to one of 6 treatments. The treatments differed from one another along two dimensions. The first dimension consisted of varying the first 5 images shown to the worker. This was used to test the effects of *in-task* priming.

The second dimension concerned disclosure of a (fictitious) organization, purportedly funding

work as part of a research study. Depending on the treatment, one of two funding agencies was presented, or no indication was made.

The 5 test images, were chosen carefully with two ideals in mind. First, we chose images that we judged would generate a diverse vocabulary of labels, such that the effects of priming could be detected. In other words, sparse images with a single object in the foreground were not considered good candidates, since they would be less likely to elicit labels that varied from one worker and one priming treatment to the next.

Second, we chose images which would produce labels belonging to two broad concepts, which would serve as the targets of our priming: food and culture. This created the opportunity to attempt to prime workers in a way that would bias them toward emitting food-related or culture-related labels.

Under these considerations, we chose the images shown in Fig.

We chose the three sets of priming images to that end. The first set, was chosen to be for the treatment

*culture*, and *food*.

## Results

## References

- [1] Dana Chandler and Adam Kapelner. Breaking monotony with meaning: Motivation in crowdsourcing markets. *Journal of Economic Behavior & Organization*, 90:123–133, 2013.

1



2



3



4



5



1



2



3



4



5





1



2



3



4



5



6

